

# Context-Aware Image Inpainting for Automatic Object Removal

---

Nicholas Russell

# What is Image Inpainting?

## Overview:

- Realistically filling or reconstructing parts of an image that are damaged or missing.

## Applications:

- Image Editing
- Image Reconstruction
- Painting or Image Restoration
- Object Removal

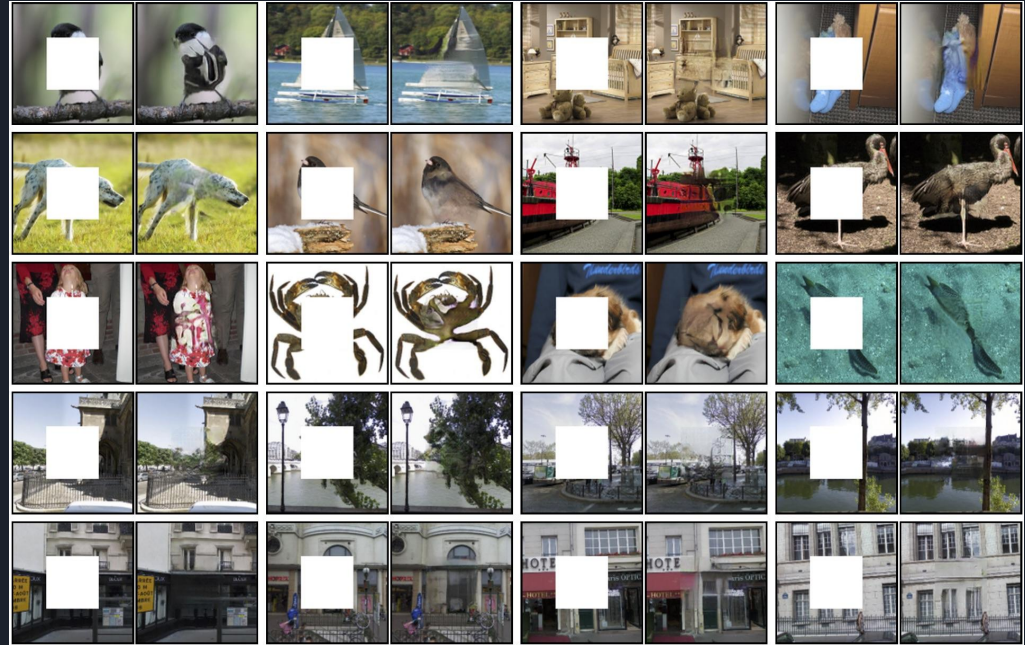
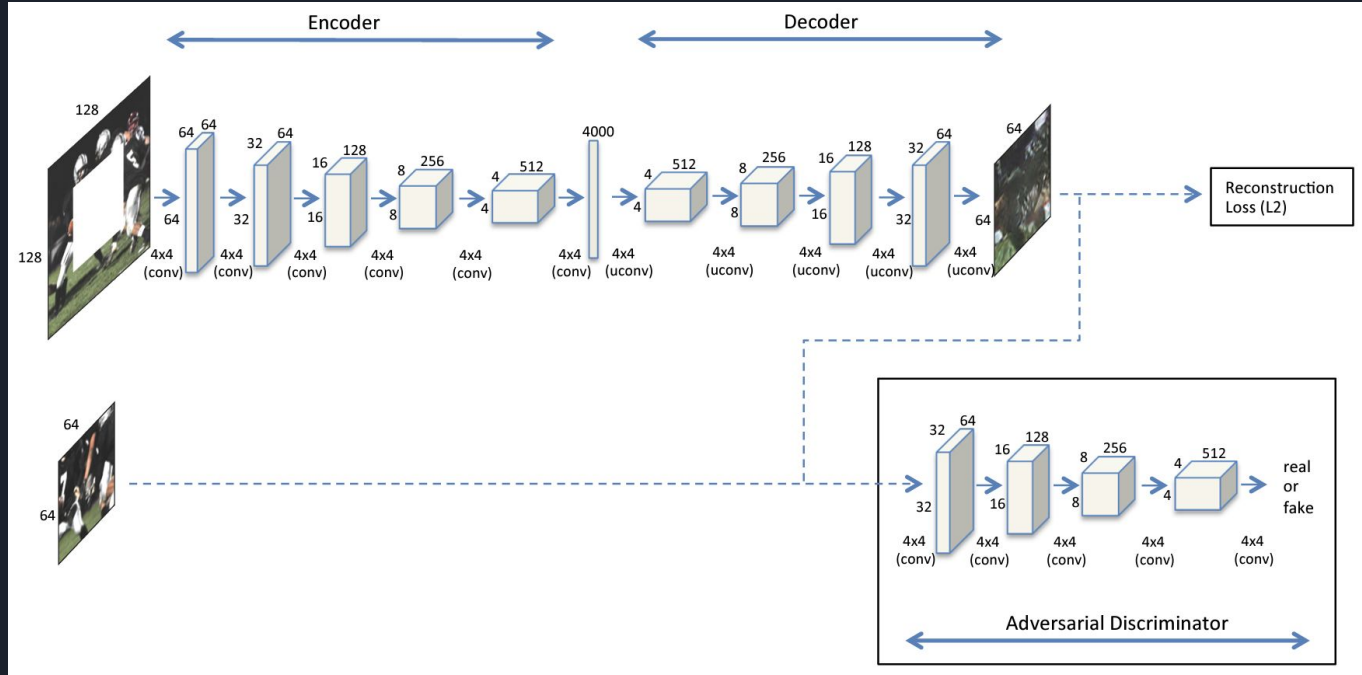


Image from [https://www.cs.cmu.edu/~dpathak/context\\_encoder/](https://www.cs.cmu.edu/~dpathak/context_encoder/)

# Image Inpainting Examples



# Context Encoders




Context Encoders Architecture [1]

Loss:  $\mathcal{L} = \lambda_{adv}\mathcal{L}_{adv} + \lambda_{rec}\mathcal{L}_{rec}$  with BCE adversarial and L2 reconstruction losses.



# Context Encoders Example



- Reconstruction loss alone results in good structure, but blurry results.
- Solution: Use a decoupled loss with an adversarial factor to exploit the structure of the reconstruction loss and the sharpness of the adversarial loss.



# Methodology

- Experiment with reconstruction metrics and implement different joint reconstruction losses to achieve better visual results.
- Modify the model for automatic object removal:
  - Modify the baseline Context Encoders model to support evaluation with arbitrary masks.
  - Integrate YOLOv8, an object detection model, with Context Encoders to generate masks to remove objects from the scene.

# Dataset

- We used the MiniPlaces dataset, which is a subset of MIT's Places dataset that only contains 100,000 128x128 images from 100 scene categories.





# Metrics

- L1

$$L1 = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$$

- L2

$$L2 = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- PSNR


$$PSNR = 10 \cdot \log_{10} \left( \frac{(MAX_I)^2}{\sqrt{MSE}} \right)$$

- SSIM

$$SSIM(x, y) = \frac{(2\mu_x\mu_y + c_1)(2\sigma_{xy} + c_2)}{(\mu_x^2 + \mu_y^2 + c_1)(\sigma_x^2 + \sigma_y^2 + c_2)}$$



# Initial Results & Motivation for Proposed Solution




Input



Masked





Base

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

# Baseline vs L2 Adversarial + L1 Reconstruction

Base



$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$



# Baseline vs L2 Adversarial + L1 Reconstruction

Base

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$



# Baseline vs L2 Adversarial + L1 Reconstruction

Base

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$







# Experiments


- What we want:
  - Bring back sharp edges, but create smoother results with less noise than some of the patches created by the baseline Context Encoders model.
- Proposed Solution:
  - Implement a Joint Reconstruction Loss specifically designed for visual appearance

Reconstruction Loss:  $\lambda_{rec1}\mathcal{L}_{rec1} + \lambda_{rec2}\mathcal{L}_{rec2}$

- Use a factor of SSIM (or variations of SSIM):

Loss:  $\mathcal{L} = \lambda_{adv}\mathcal{L}_{adv} + \lambda_{rec}(\lambda_{L1}\mathcal{L}_{L1} + \lambda_{SSIM}(1 - \mathcal{L}_{SSIM}))$

Base



L1: 0.0774  
L2: 0.0386  
PSNR: 14.1375  
SSIM: 0.7189

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0698  
L2: 0.0347  
PSNR: 14.5935  
SSIM: 0.7403



L1: 0.0734  
L2: 0.0396  
PSNR: 14.0252  
SSIM: 0.7287

L1: 0.0844  
L2: 0.0552  
PSNR: 12.5826  
SSIM: 0.7232

L1: 0.0744  
L2: 0.0400  
PSNR: 13.9811  
SSIM: 0.7345

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$

Base



L1: 0.0551  
L2: 0.0199  
PSNR: 17.0017  
SSIM: 0.7063

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$



L1: 0.0511  
L2: 0.0180  
PSNR: 17.4431  
SSIM: 0.7164

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$



L1: 0.0562  
L2: 0.0221  
PSNR: 16.5612  
SSIM: 0.7139

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$



L1: 0.0560  
L2: 0.0213  
PSNR: 16.7069  
SSIM: 0.7128



L1: 0.0661  
L2: 0.0298  
PSNR: 15.2601  
SSIM: 0.7060



Base



L1: 0.0509  
L2: 0.0209  
PSNR: 16.8047  
SSIM: 0.7701

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0415  
L2: 0.0181  
PSNR: 17.4288  
SSIM: 0.8060



L1: 0.0522  
L2: 0.0290  
PSNR: 15.3750  
SSIM: 0.7975

L1: 0.0507  
L2: 0.0224  
PSNR: 16.4965  
SSIM: 0.7920



L1: 0.0448  
L2: 0.0205  
PSNR: 16.8840  
SSIM: 0.8065

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$

Base



L1: 0.0691  
L2: 0.0296  
PSNR: 15.2889  
SSIM: 0.7111

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0648  
L2: 0.0268  
PSNR: 15.7234  
SSIM: 0.7194



L1: 0.0640  
L2: 0.0279  
PSNR: 15.5403  
SSIM: 0.7256

L1: 0.0789  
L2: 0.0415  
PSNR: 13.8219  
SSIM: 0.7187

L1: 0.0669  
L2: 0.0289  
PSNR: 15.3979  
SSIM: 0.7197


$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$



Base




L1: 0.0652  
L2: 0.0277  
PSNR: 15.5769  
SSIM: 0.7183

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0596  
L2: 0.0244  
PSNR: 16.1256  
SSIM: 0.7237



L1: 0.0613  
L2: 0.0277  
PSNR: 15.5784  
SSIM: 0.7278



L1: 0.0662  
L2: 0.0313  
PSNR: 15.0510  
SSIM: 0.7267



L1: 0.0607  
L2: 0.0260  
PSNR: 15.8559  
SSIM: 0.7234

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$

Base



L1: 0.0740  
L2: 0.0358  
PSNR: 14.4571  
SSIM: 0.7388

# Results

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$



L1: 0.0688  
L2: 0.0389  
PSNR: 14.1060  
SSIM: 0.7421



$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$



L1: 0.0724  
L2: 0.0417  
PSNR: 13.8002  
SSIM: 0.7382

L1: 0.0811  
L2: 0.0553  
PSNR: 12.5762  
SSIM: 0.7232


$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$



L1: 0.0734  
L2: 0.0464  
PSNR: 13.3350  
SSIM: 0.7479



Base



L1: 0.0783  
L2: 0.0360  
PSNR: 14.4411  
SSIM: 0.7140

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0846  
L2: 0.0468  
PSNR: 13.2932  
SSIM: 0.7185



L1: 0.0941  
L2: 0.0572  
PSNR: 12.4256  
SSIM: 0.7129

L1: 0.1020  
L2: 0.0716  
PSNR: 11.4505  
SSIM: 0.7136

L1: 0.0859  
L2: 0.0474  
PSNR: 13.2443  
SSIM: 0.7118

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$

$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$

$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$

Base



L1: 0.0954  
L2: 0.0521  
PSNR: 12.8282  
SSIM: 0.7108

# Results

$$\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$$

L1: 0.0918  
L2: 0.0516  
PSNR: 12.8775  
SSIM: 0.7213



L1: 0.0962  
L2: 0.0609  
PSNR: 12.1527  
SSIM: 0.7269

L1: 0.1137  
L2: 0.0873  
PSNR: 10.5899  
SSIM: 0.7164



very high SSIM in combination with adversarial loss tends to give strange results or artifacts

$$\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$$



$$\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$$



$$\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$$

L1: 0.0915  
L2: 0.0543  
PSNR: 12.6535  
SSIM: 0.7365

# Evaluation

- Each model was trained for 40 epochs.
- Models are evaluated on the MiniPlaces validation dataset containing 10,000 128x128 images from 100 different scene categories.
- All methods, except the second, use BCE adversarial loss.

| Model  | Mean L1 Loss  | Mean L2 Loss  | Mean PSNR      | Mean SSIM     |
|--|---------------|---------------|----------------|---------------|
| Base Context Encoders  | 0.0653        | 0.0319        | 15.7722        | 0.7385        |
| $\mathcal{L}_{adv} = \mathcal{L}_{L2}, \mathcal{L}_{rec} = \mathcal{L}_{L1}$ | <b>0.0619</b> | <b>0.0319</b> | <b>15.9358</b> | <b>0.7538</b> |
| $\mathcal{L}_{rec} = 0.3\mathcal{L}_{L1} + 0.7\mathcal{L}_{SSIM}$            | 0.0651        | 0.0366        | 15.3956        | 0.7504        |
| $\mathcal{L}_{rec} = 0.5\mathcal{L}_{L1} + 0.5\mathcal{L}_{SSIM}$            | 0.0645        | 0.0363        | 15.4440        | 0.7512        |
| $\mathcal{L}_{rec} = 0.005\mathcal{L}_{L1} + 0.995\mathcal{L}_{SSIM}$        | 0.0740        | 0.0485        | 14.3107        | 0.7441        |



# Object Removal – Mask Extraction


Input



Predictions



Mask




Extracting Mask of People and Cars

# Automatic Object Removal Examples



- Note: The model is trained 12.5x less than in the Context Encoders paper, so the inpainting results are not perfect and object outlines can still be seen.

# Object Removal Limitations



- Problem: Removal with arbitrary masks sometimes gives strange artifacts or inpaints a near solid color (especially on small objects).
- Potential Solution: Implement random masking when training the model.



# Conclusion

- Modified the baseline Context Encoders model to support evaluation with arbitrary shaped masks.
- Implemented the below joint reconstruction loss to significantly improve visual results of the Context Encoders model.

$$\text{Loss: } \mathcal{L} = \lambda_{adv}\mathcal{L}_{adv} + \lambda_{rec}(\lambda_{L1}\mathcal{L}_{L1} + \lambda_{SSIM}(1 - \mathcal{L}_{SSIM}))$$

- Integrated image inpainting with object detection and segmentation to automatically remove objects from an image.



# References

- [1] Deepak Pathak, Philipp Krähenbühl, Jeff Donahue, Trevor Darrell, and Alexei Efros. “Context Encoders: Feature Learning by Inpainting”. In: *Computer Vision and Pattern Recognition (CVPR)*. 2016.
- [2] Hang Zhao, Orazio Gallo, Iuri Frosio, and Jan Kautz. “Loss Functions for Image Restoration With Neural Networks”. In: *IEEE Transactions on Computational Imaging* 3.1 (2017), pp. 47–57.
- [3] Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba. “Places: A 10 million Image Database for Scene Recognition”. In: *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2017).